

# **Reinforcement Learning for Improving the Performance of Recovery Algorithm in Noisy Group Testing**

Seyedeh Nahid Esmati, Arghamitra Talukder, and Ya Mei

Department of Electrical and Computer Engineering

Texas A&M University

Professor: Dr. Dileep Kalathil

May 5, 2021

# Tasks Done by Each Team Member

## Seyedeh Nahid Esmati:

- Implementing advantage actor-critic policy gradient with four different notions of states, and four different notions of rewards as stated in the slides.
- For the baseline work on GT and LDPC: Implementing the BP recovery algorithm (without RL) with flooding and maximum-residual node scheduling policies for both GT and LDPC, in order to reproduce the benchmarks.
- For the baseline paper on LDPC: Implementing Q-learning with direct update rule (no function approximation)+clustering+quantization.

## Arghamitra Talkuder:

- Implementing DQN with multiple agents (NNs) to learn multiple action-value functions in parallel (one agent for each BP iteration) for the state and reward definitions stated in the slides.
- Monte-Carlo simulation to reproduce the benchmark results for the BP algorithm with maximal-residual node scheduling policy for GT.


## Ya Mei:

- Implementing DQN with a single agent (NN) to learn a universal action-value function (one agent for all BP iterations) with two different notions of states, and two different notions of rewards as stated in the slides.
- Monte-Carlo simulation to reproduce the benchmark results for the BP algorithm with flooding for GT.

# Group Testing


Design testing algorithms by using (binary) tests to identify all defective items among a much larger group of items.

## Individual Testing



# tests (per item) = 1

## Dorfman's Testing (1943)



# tests (per item) =  $\frac{5}{6}$

How to minimize # tests for an unknown config. of def. items?

# Applications of Group Testing

Subject Journals Books Major Reference Works Partner With Us Open Access About Us

**Series on Applied Mathematics: Volume 18**  
**Pooling Designs and Nonadaptive Group Testing**  
Important Tools for DNA Sequencing

<https://doi.org/10.1142/9122> | June 2006  
Pages: 248  
By (author): Ding Zhu Du (University of Texas at Dallas, USA & Xiamen University, China) and Frank K. Hwang (National Chiao Tung University, Taiwan, ROC)

Genetic Mapping and DNA Sequencing pp 133-194 | Cite as

**A Comparative Survey of Non-Adaptive Pooling Designs**

Authors Authors and affiliations

D. J. Building, W. J. Bruno, D. C. Torrey, C. Knell

Subject Journals Books Major Reference Works Partner With Us Open Access About Us

**Series on Applied Mathematics: Volume 3**  
**Combinatorial Group Testing and Its Applications**

<https://doi.org/10.1142/9136> | November 1993  
Pages: 264  
By (author): D-Z Du (Univ Minnesota) and F K Hwang (AT&T Bell Labs)

- Public health
- Cybersecurity
- Bloom filters
- Data science
- Information theory

International Conference on Applied Cryptography and Network Security  
ACNS 2005: Applied Cryptography and Network Security pp 206-221 | Cite as

**Indexing Information for Data Forensics**

Authors Authors and affiliations

Michael T. Goodrich, Mikhail J. Atallah, Roberto Tamassia

## Born Again Group Testing: Multiaccess Communications

Invited Paper

JACK K. WOLF, FELLOW, IEEE

THE CONVERSATION

COVID-19 CORONAVIRUS TEST  
POSITIVE ☐ NEGATIVE ☐

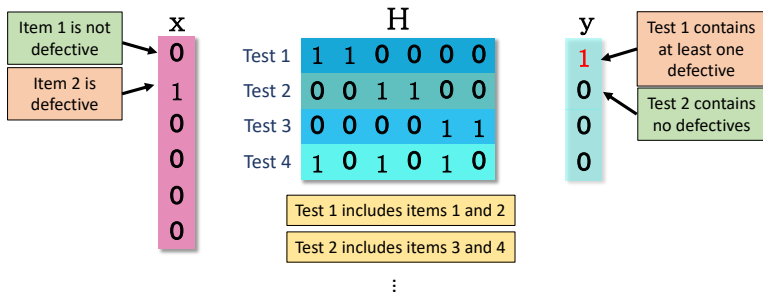
Group testing for coronavirus – called pooled testing – could be the fastest and cheapest way to increase screening nationwide

### The mathematical strategy that could transform coronavirus testing

Four charts show how pooling samples from many people can save time or resources.


# Problem Setup

- Population of  $n$  items (labeled  $1, \dots, n$ )
- $k \ll n$  are defective (i.e., probability of defective =  $q = \frac{k}{n}$ )
- $\mathbf{x} = [x_1, \dots, x_n]$ : binary vector representing status of the items
- $H \in \{0, 1\}^{m \times n}$ : binary matrix representing the testing matrix




# Noiseless Setting vs. Noisy Setting

## Noiseless:



## Noisy:



**BSC** ( $\epsilon_0 = \epsilon_1 = \delta$ ); **Z-channel** ( $\epsilon_0 = 0$ ); **reverse Z-channel** ( $\epsilon_1 = 0$ ).

# Exact Recovery vs. Partial Recovery

**Exact Recovery:** (no false negatives/positives)

- **Requirement:**  $\hat{x} = x$ ;
- **Success probability:**  $\Pr(\hat{x} = x)$ .

**Partial Recovery:** ( $\leq d_{\max}$  false neg.'s and  $\leq d_{\max}$  false pos.'s)

- **Requirement:**  $d(x, \hat{x}) \leq d_{\max}$ , where

$$d(x, \hat{x}) = \max\{|\text{supp}(x) \setminus \text{supp}(\hat{x})|, |\text{supp}(\hat{x}) \setminus \text{supp}(x)|\}$$

and

$\text{supp}(x)$  = index set of all nonzero components of  $x$ .

- **Success probability:**  $\Pr(d(x, \hat{x}) \leq d_{\max})$ .

# Recovery Algorithm: Belief Propagation (BP)

Check-to-variable messages:  $m_{c \rightarrow v}^{(l)}(u_v)$

If  $z_c = 0$ :

$$\propto \begin{cases} \delta & u_v = 1 \\ \delta + (1 - 2\delta) \prod_{v' \in \mathcal{N}(c) \setminus \{v\}} m_{v' \rightarrow c}^{(l-1)}(0) & u_v = 0 \end{cases}$$

If  $z_c = 1$ :

$$\propto \begin{cases} 1 - \delta & u_v = 1 \\ 1 - \delta - (1 - 2\delta) \prod_{v' \in \mathcal{N}(c) \setminus \{v\}} m_{v' \rightarrow c}^{(l-1)}(0) & u_v = 0 \end{cases}$$

Variable-to-check messages:  $m_{v \rightarrow c}^{(l)}(u_v)$

$$\propto \begin{cases} q \prod_{c' \in \mathcal{N}(v) \setminus \{c\}} m_{c' \rightarrow v}^{(l)}(u_v) & u_v = 1 \\ (1 - q) \prod_{c' \in \mathcal{N}(v) \setminus \{c\}} m_{c' \rightarrow v}^{(l)}(u_v) & u_v = 0 \end{cases}$$

Variable Nodes  
(VNs)

$v_1$

$v_2$

$v_3$

$v_4$

$v_5$

$v_6$

Items

Check Nodes  
(CNs)


$c_1$

$c_2$

$c_3$

$c_4$

Tests





# Recovery Algorithm: Belief Propagation (BP)

Log Likelihood Ratio (LLR):  $L_v^{(l)}$

$$= \begin{cases} \ln \frac{q}{1-q} & l = 0 \\ \ln \frac{q}{1-q} + \sum_{c \in \mathcal{N}(v)} \ln \frac{m_{c \rightarrow v}^{(l)}(1)}{m_{c \rightarrow v}^{(l)}(0)} & l \geq 1 \end{cases}$$

Decoding Rule:

- If  $L_v^{(l)} > 0$ :  $\hat{x}_v^{(l)} = 1$
- If  $L_v^{(l)} = 0$ :  $\hat{x}_v^{(l)} = 1$  or  $0$  w.p.  $\frac{1}{2}$
- If  $L_v^{(l)} < 0$ :  $\hat{x}_v^{(l)} = 0$

Variable Nodes  
(VNs)

$v_1$

$v_2$

$v_3$

$v_4$

$v_5$

$v_6$

Items

Check Nodes  
(CNs)


$c_1$

$c_2$

$c_3$

$c_4$


Tests



# Flooding vs. Node Scheduling

**Flooding:** In each iteration, all CNs send messages to their neighboring VNs, and all VNs send messages to their neighboring CNs.

**Node Scheduling:** In each iteration, a single CN sends messages to its neighboring VNs, and these VNs send messages to their neighboring CNs.



**Goal:** Given testing matrix  $H$ , # defectives  $k$ , and crossover prob.  $\delta$ , find an optimal node scheduling policy that maximizes the success probability.

# Apply RL to Learn Optimal Node Scheduling Policy

- Scheduling a single check node at each iteration of the BP algo. can be viewed as a Multi-Armed Bandits (MAB) process.

- **State:**

- (1) The sum of the LLRs of the VNs connected to each CN,

$$\left[ \sum_{v \in \mathcal{N}(1)} L_v^{(l)}, \sum_{v \in \mathcal{N}(2)} L_v^{(l)}, \dots, \sum_{v \in \mathcal{N}(m)} L_v^{(l)} \right],$$

where

$$L_v^{(l)} = \begin{cases} \ln \frac{q}{1-q} & l = 0 \\ \ln \frac{q}{1-q} + \sum_{c \in \mathcal{N}(v)} \ln \frac{m_{c \rightarrow v}^{(l)}(1)}{m_{c \rightarrow v}^{(l)}(0)} & l \geq 1 \end{cases}.$$

- (2) The LLRs of the VNs connected to each CN,

$$\left[ \{L_v^{(l)}\}_{v \in \mathcal{N}(1)}, \{L_v^{(l)}\}_{v \in \mathcal{N}(2)}, \dots, \{L_v^{(l)}\}_{v \in \mathcal{N}(m)} \right].$$

- (3) A quantized version of (1) or (2);

- (4) (1), (2), or (3) + the observed noisy test results  $\hat{y}_1, \dots, \hat{y}_m$ .

# Apply RL to Learn Optimal Node Scheduling Policy

- **Reward:**

- (1) Max. change in LLR of a variable node incident to the currently-scheduled check node  $c^*$ ,

$$r_{c^*} = \max_{v \in \mathcal{N}(c^*)} |L_{c^* \rightarrow v}^{\text{new}} - L_{c^* \rightarrow v}^{\text{old}}|,$$

where

$$L_{c \rightarrow v}^{\text{new}} = \ln \frac{m_{c \rightarrow v}^{\text{new}}(1)}{m_{c \rightarrow v}^{\text{new}}(0)}, \quad L_{c \rightarrow v}^{\text{old}} = \ln \frac{m_{c \rightarrow v}^{\text{old}}(1)}{m_{c \rightarrow v}^{\text{old}}(0)}.$$

- (2) Ratio of the max. change in LLR of a variable node incident to the currently-scheduled check node  $c^*$  to the max. change in LLR of any variable node when scheduling any check node  $c$ ,

$$R = \frac{r_{c^*}}{\max_{c \in [m]} r_c}.$$

- (3) Minus the Hamming distance between the original signal  $\mathbf{x} = [x_1, x_2, \dots, x_n]$  and the currently-estimated signal  $\hat{\mathbf{x}}^{(l)}$ ,

$$\hat{\mathbf{x}}^{(l)} = \left[ \mathbb{1}_{(L_1^{(l)} > 0)}, \mathbb{1}_{(L_2^{(l)} > 0)}, \dots, \mathbb{1}_{(L_n^{(l)} > 0)} \right].$$

- (4) A linear/non-linear function of (1), (2), or (3) to further encourage higher rewards and/or discourage lower rewards.

# Case Study

## Problem parameters:

- # items  $n = 100$
- # tests  $m = 20$
- # defective items  $k = 2$
- Testing matrix  $H$ : A randomly generated  $20 \times 100$  Bernoulli matrix with probability  $1 = \ln(2)/2 \approx 0.35$
- Crossover probability of noisy channel  $\epsilon = 0.05$
- # BP iterations: flooding = 10, node scheduling = 100

## Benchmarks:

	Flooding	Maximum-Residual Node Scheduling
Success Probability (%) (Unknown $k$ )	32.99	51.22
Success Probability (%) (Known $k$ )	50.00	61.34

# Policy Gradient (Advantage Actor-Critic)

**State:** Sum of the LLRs of the VNs connected to each CN,

$$\left[ \sum_{v \in \mathcal{N}(1)} L_v^{(l)}, \sum_{v \in \mathcal{N}(2)} L_v^{(l)}, \dots, \sum_{v \in \mathcal{N}(m)} L_v^{(l)} \right]$$

where

$$L_v^{(l)} = \begin{cases} \ln \frac{q}{1-q} & l = 0 \\ \ln \frac{q}{1-q} + \sum_{c \in \mathcal{N}(v)} \ln \frac{m_{c \rightarrow v}^{(l)}(1)}{m_{c \rightarrow v}^{(l)}(0)} & l \geq 1 \end{cases}.$$

**Reward:** A non-linear function of the ratio  $R$ ,

$$10 \tanh(R - 0.5)$$

where

$$R = \frac{\max_{v \in \mathcal{N}(c^*)} |L_{c^* \rightarrow v}^{\text{new}} - L_{c^* \rightarrow v}^{\text{old}}|}{\max_{c \in [m]} \max_{v \in \mathcal{N}(c)} |L_{c \rightarrow v}^{\text{new}} - L_{c \rightarrow v}^{\text{old}}|},$$
$$L_{c \rightarrow v}^{\text{new}} = \ln \frac{m_{c \rightarrow v}^{\text{new}}(1)}{m_{c \rightarrow v}^{\text{new}}(0)}, \quad L_{c \rightarrow v}^{\text{old}} = \ln \frac{m_{c \rightarrow v}^{\text{old}}(1)}{m_{c \rightarrow v}^{\text{old}}(0)}.$$

# Results: Policy Gradient (Advantage Actor-Critic)

## Learning parameters:

- Learning rate  $\alpha = 0.001$
- Discount factor  $\gamma = 0.99$

## Actor and critic NNs:

- # hidden layers: 2
- # nodes per layer: 60



	Flooding	Max-Res NS	RL-based Node Scheduling					
			500 eps	1000 eps	1500 eps	2000 eps	2500 eps	3000 eps
Success Probability (%) (Unknown $k$ )	32.99	51.22	35.4	45.5	47.1	48.3	48.9	49.2
Success Probability (%) (Known $k$ )	50.00	61.34	47.2	56.4	57.6	58.1	58.5	59.0

# DQN with Multiple NNs

(One NN for Each BP Iterations)

## Following the sequential steps as code: Declaring agent.

**Goal:** To train  $I_{\max}$  NNs, each as a nonlinear approx. of the action-value function  $Q^{(I)}(\text{state}, \text{action})$  for the BP iteration  $1 \leq I \leq I_{\max}$ .  
Individual agent for each iteration

For state (1) and reward (1):

- # hidden layers: 2
- # nodes per layer: {32, 64}
- Batch size: 64
- Replay buffer size:  $1e5$
- Learning rate  $\alpha$ :  $1e-4$
- Discount factor  $\gamma$ : 0.99
- Decay rate ( $\epsilon$ -greedy): 0.999
- Iteration per episode: 101



# DQN with Multiple NNs

## (One NN for Each BP Iterations)



### First observation

**State:** The sum of the LLRs of the VNs connected to each CN,

$$\left[ \sum_{v \in \mathcal{N}(1)} L_v^{(l)}, \sum_{v \in \mathcal{N}(2)} L_v^{(l)}, \dots, \sum_{v \in \mathcal{N}(m)} L_v^{(l)} \right],$$

where

$$L_v^{(l)} = \begin{cases} \ln \frac{q}{1-q} & l = 0 \\ \ln \frac{q}{1-q} + \sum_{c \in \mathcal{N}(v)} \ln \frac{m_{c \rightarrow v}^{(l)}(1)}{m_{c \rightarrow v}^{(l)}(0)} & l \geq 1 \end{cases}$$




# DQN with Multiple NNs

(One NN for Each BP Iterations)

## First iteration: Choose action

$\epsilon$ -greedy: For first iteration selecting a random node. Action space is same as state space

## First iteration: New observation based on chosen action



# DQN with Multiple NNs

(One NN for Each BP Iterations)


## First iteration: Reward

**Reward:** The max. change in LLR of a variable node incident to the currently-scheduled check node  $c^*$ ,

$$\max_{v \in \mathcal{N}(c^*)} |L_{c^* \rightarrow v}^{\text{new}} - L_{c^* \rightarrow v}^{\text{old}}|,$$

where

$$L_{c \rightarrow v}^{\text{new}} = \ln \frac{m_{c \rightarrow v}^{\text{new}}(1)}{m_{c \rightarrow v}^{\text{new}}(0)}, \quad L_{c \rightarrow v}^{\text{old}} = \ln \frac{m_{c \rightarrow v}^{\text{old}}(1)}{m_{c \rightarrow v}^{\text{old}}(0)}.$$



# DQN with Multiple NNs

## (One NN for Each BP Iterations)

### First iteration: After reward

- Store in replay buffer
- Learn: loss, backpropagation, step




Table: Success probability for different episodes

Success Probability	500	1500	2500	3000
% for unknown k	39.2	41.3	46.9	48.7
% for known k	50.2	53.5	58.3	60.3

- We see a gradual increase in success probability as the number of episodes increases.
- Reward remains almost same, irrespective of episodes (around 55); and a gradual learning can be noticed from episode 0-500.

# DQN with a Single NN

## (One NN for All BP Iterations)

State:

- (1) The LLRs of the VNs connected to each CN,

$$\left[ \{L_v^{(l)}\}_{v \in \mathcal{N}(1)}, \{L_v^{(l)}\}_{v \in \mathcal{N}(2)}, \dots, \{L_v^{(l)}\}_{v \in \mathcal{N}(m)} \right].$$

- (2) The sum of the LLRs of the VNs connected to each CN,

$$\left[ \sum_{v \in \mathcal{N}(1)} L_v^{(l)}, \sum_{v \in \mathcal{N}(2)} L_v^{(l)}, \dots, \sum_{v \in \mathcal{N}(m)} L_v^{(l)} \right].$$

Reward:

- (1) Minus the Hamming distance between the original signal  $\mathbf{x} = [x_1, x_2, \dots, x_n]$  and the currently-estimated signal  $\hat{\mathbf{x}}^{(l)}$ ,

$$\hat{\mathbf{x}}^{(l)} = \left[ \mathbb{1}_{(L_1^{(l)} > 0)}, \mathbb{1}_{(L_2^{(l)} > 0)}, \dots, \mathbb{1}_{(L_n^{(l)} > 0)} \right].$$

- (2) The max. change in LLR of a variable node incident to the currently-scheduled check node  $\mathbf{c}^*$ ,

$$\max_{v \in \mathcal{N}(\mathbf{c}^*)} |L_{\mathbf{c}^* \rightarrow v}^{\text{new}} - L_{\mathbf{c}^* \rightarrow v}^{\text{old}}|.$$


# DQN with a Single NN

## (One NN for All BP Iterations)

**Goal:** To train a neural network (NN) as a nonlinear approx. of the action-value function  $Q(\text{state}, \text{action})$ .

For state (1) and reward (1):

- # hidden layers: {1, 2, 3, 4, 5}
- # nodes per layer (3 layers): {1000, 500, 100}
- Batch size (exp. replay): {4, 16, 32}
- Update freq. (target net): {4, 8, 16}
- Learning rate  $\alpha$ : {0.01, 0.005, 0.001, 0.0001}
- Discount factor  $\gamma$ : {0.9, 0.95, 0.99}
- Decay rate ( $\epsilon$ -greedy): {0.99, 0.999, 0.9999, 0.99999}



Training: Not successful!  
(Success probability = 0%)

# Summary

**Goal:** To find a node scheduling (NS) policy that maximizes the success probability of the BP recovery algo. for noisy group testing.

Our RL-based NS policies (actor-critic and DQN + multiple NNs):

- outperform the flooding approach in terms of success probability.
- perform closely to the maximum-residual NS policy in terms of success probability, and are less computationally complex.

## Future Work

- Applying variance reducing methods e.g. Averaged-DQN, which is an extension to the DQN based on averaging previously learned Q-values estimates.
- Improving PG by applying a more directed exploration strategy that promotes exploration of under-appreciated rewards.
- Applying other policy gradient algorithms e.g., Natural Policy Gradient (NPG), Proximal Policy Optimization (PPO).
- Implementing DQN + multiple NNs with target network to improve the performance.
- Different variations of the problem:
  - BSC with different cross-over probabilities
  - Z-channel and reverse Z-channel
  - Partial recovery guarantee